

General Neural Network Questions

1. **Why sigmoid in the output layer for binary classification?**
Sigmoid outputs probabilities (0 to 1) for binary classification.
2. **What does `input_dim=X.shape[1]` mean in the Dense layer?**
It sets the input layer to match the number of features in X.
3. **Purpose of StandardScaler in preprocessing?**
Scales features to zero mean and unit variance for faster convergence.
4. **What does `batch_size=32` mean in `model.fit`?**
Processes 32 samples per gradient update, balancing speed and stability.
5. **Why binary_crossentropy as loss? What for multiclass?**
Measures binary classification error; use categorical_crossentropy for multiclass.
6. **Role of adam optimizer?**
Adapts learning rate for efficient weight updates.
7. **Why relu in hidden layers?**
Adds non-linearity, avoids vanishing gradients, and is computationally efficient.
8. **Why split data with train_test_split?**
Trains on one set, tests on unseen data to evaluate generalization.

Convolutional Neural Network (CNN) Questions

9. **Purpose of Conv2D layer?**
Extracts spatial features like edges from images using filters.
10. **Why use MaxPooling2D after Conv2D?**
Reduces dimensions, lowers computation, and prevents overfitting.
11. **What does Flatten layer do?**
Converts 2D feature maps to 1D for dense layers.
12. **Significance of `rescale=1./255` in ImageDataGenerator?**
Normalizes pixel values to [0, 1] for stable training.
13. **Why input shape (64, 64, 3) in CNN?**
Represents 64x64 RGB images (height, width, 3 channels).
14. **Purpose of Dropout(0.5) in custom CNN?**
Randomly drops 50% of neurons to prevent overfitting.
15. **Why more layers in custom CNN?**
Learns complex features for better pattern recognition.

16. Changes for multiclass classification in CNN?

Use softmax in output layer and categorical_crossentropy loss.

Transfer Learning Questions

17. What is transfer learning, and why used?

Uses pre-trained model features to save time and improve performance.

18. Why set base model layers to trainable = False?

Freezes pre-trained layers to retain learned features.

19. Role of GlobalAveragePooling2D?

Averages feature maps to a vector for classification.

20. Why input shape (224, 224, 3) for VGG16?

Matches pre-trained model's expected 224x224 RGB input.

21. Difference between categorical_crossentropy and binary_crossentropy?

categorical for multiclass, binary for two-class problems.

22. How does Adam differ from SGD?

Adam uses adaptive learning rates for faster convergence.

Recurrent Neural Network (RNN, LSTM, GRU) Questions

23. Why LSTM for stock price prediction?

Captures long-term dependencies in time-series data.

24. What is sequence_length=60 in create_sequences?

Uses 60 past time steps to predict the next value.

25. Why use MinMaxScaler for stock prices?

Scales data to [0, 1] for stable training.

26. Difference between SimpleRNN, LSTM, GRU?

SimpleRNN has vanishing gradients; LSTM and GRU handle long-term dependencies, with GRU simpler.

27. Why mean_squared_error for stock prediction?

Minimizes error for continuous regression predictions.

28. How does Dense(1) produce predictions in LSTM?

Outputs one value for the next stock price.

Next Word Prediction Questions

29. Role of Embedding layer?

Converts words to dense vectors for semantic representation.

- 30. **Why softmax in output layer?**
Outputs probabilities over vocabulary for next word.
- 31. **What does to_categorical do?**
Converts word indices to one-hot encoded vectors.
- 32. **Why calculate max_seq_len?**
Pads sequences to fixed length for consistent input.
- 33. **How does LSTM help in next word prediction?**
Captures contextual dependencies in text sequences.

Autoencoder Questions

- 34. **Purpose of autoencoder in MNIST code?**
Learns compressed data representations for reconstruction.
- 35. **Why input/output shape (784,) in autoencoders?**
Represents flattened 28x28 MNIST images.
- 36. **Difference between vanilla and deep autoencoder?**
Vanilla has one encoding/decoding layer; deep has multiple.
- 37. **Why binary_crossentropy in autoencoders?**
Measures pixel-wise error for normalized images.
- 38. **Role of UpSampling2D in convolutional autoencoder?**
Increases feature map size for image reconstruction.
- 39. **How does convolutional autoencoder differ?**
Uses convolutional layers for spatial data, unlike dense layers.
- 40. **Why normalize MNIST to [0, 1]?**
Stabilizes training and matches loss function.

Q-Learning Questions

- 41. **Purpose of Q-matrix?**
Stores expected rewards for state-action pairs.
- 42. **What does gamma represent?**
Balances immediate vs. future rewards (0.8).
- 43. **Why -1 for invalid actions in R?**
Prevents agent from choosing invalid actions.
- 44. **Significance of Bellman equation?**
Updates Q-values with immediate and future rewards.

45. Why normalize Q-matrix?

Scales values to [0, 100] for interpretability.

46. What happens at terminal states (2 or 3)?

Episode ends as goals are reached.

47. How does learning rate (α) affect Q-learning?

Controls how much new rewards update Q-values.

General Deep Learning Concepts

48. Supervised vs. unsupervised learning in notebook?

Supervised (CNNs) uses labels; unsupervised (autoencoders) learns from unlabeled data.

49. Why use epochs? Effect of number?

Epochs repeat dataset passes; too few underfit, too many overfit.

50. What is overfitting? How do dropout/augmentation help?

Overfitting is poor generalization; dropout deactivates neurons, augmentation adds data variety.